

Introduction to Machine Learning



We will be starting shortly...

Introduction to Machine Learning



Module 1 Live Session

Welcome!



JOSIAH WANG

Module 0, 1
Optional lab exercises
Website/Support



ANTOINE CULLY

Modules 2, 6, 7
Coursework 1



MAREK REI

Modules 3, 4, 5
Coursework 2

CONTACT US

Ed for shared questions

Email for private issues

TEACHING ASSISTANTS

Adam Marcus

Akis Kefalas

Athanasios Vlontzos

Carles Balsells Rodas

Christina Karakosta

Ekin Ozturk

Fabrizio Russo

Hadrien Reynaud

Harry Coppock

Konstantinos Barmpas

Kyriaki-Margarita Bintsi

Najla Al Futaisi

Padmanaba Srinivasan

Rolandos-Alexandros Potamias

Turkay Kart

Today

- Q&A
 - You ask questions
 - In chat
 - Raise hands
 - We answer
- Some discussions and quizzes

Administrative questions?

Week 3 Materials (Module 2)

Released today!


~~Exam~~ TRA rubrics

3 out of 3 questions, 90 minutes

Coursework

Coursework 1

- 40% of coursework (12% of module grade)
- Implement decision trees from scratch
 - Python Standard Library
 - NumPy
 - Matplotlib



- To be released on Monday morning
 - Specs and code will be on CATE and/or Scientia
- Form your groups of 4 people
 - Can change until you submit coursework 1
- Please contribute to your group!
 - We have the right to scale your marks for problem cases!

Q: Coursework group size < 4 ?

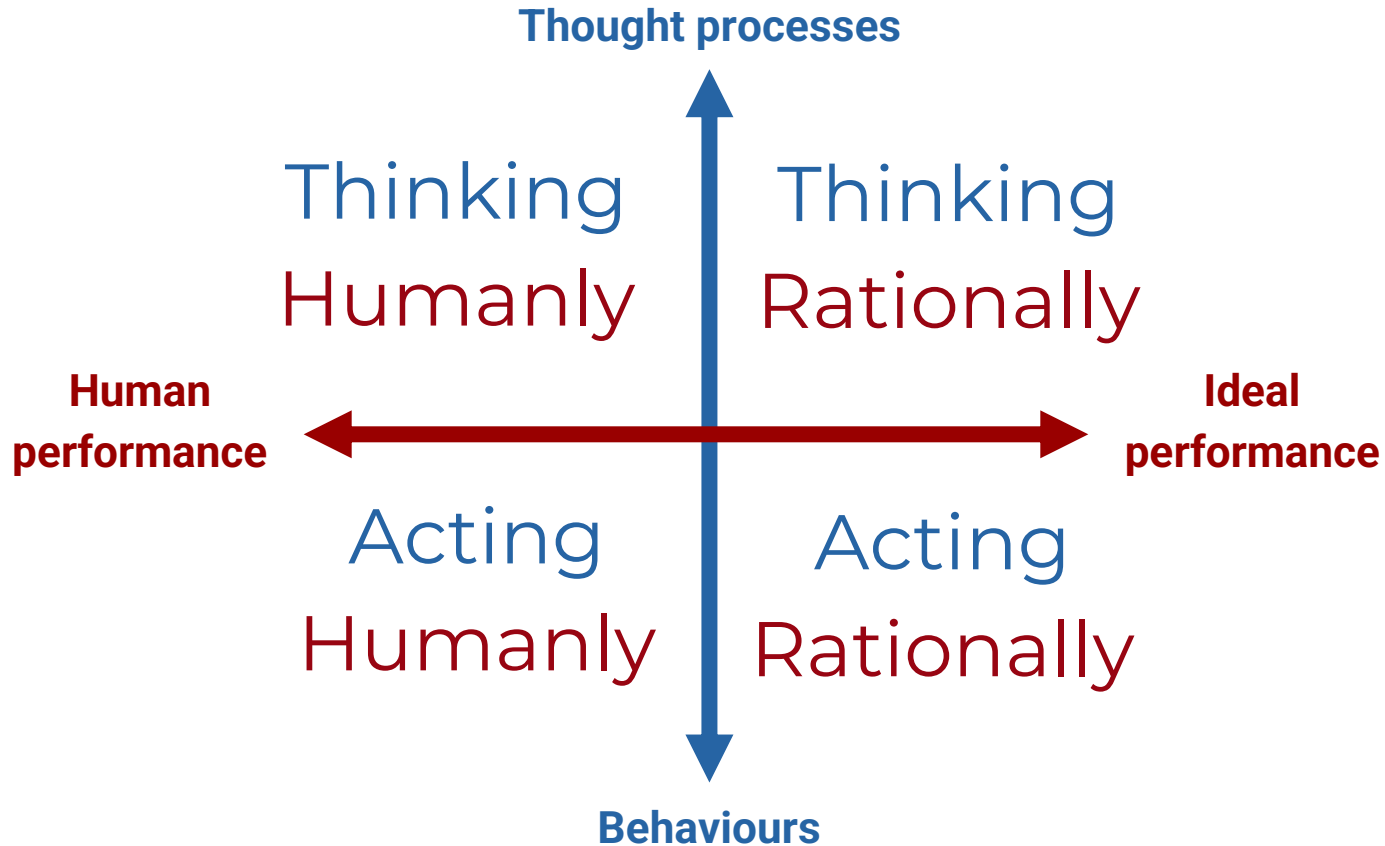
No

What is Artificial Intelligence?

How does it relate to Machine Learning?

What did you come up with?





**So, what *exactly* is
Machine Learning?**

What did *you* think?



$$f(\text{CW1}, \text{CW2}, \text{Exam}) = \text{Module Grade}$$

\approx

$$h(\text{CW1}, \text{CW2}, \text{Exam} \mid \text{D}) = \text{Estimated Grade}$$

D

😊	CW1	CW2	Exam	Grade
😊	CW1	CW2	Exam	Grade
😊	CW1	CW2	Exam	Grade

😊	CW1	CW2	Exam	Grade
😊	CW1	CW2	Exam	Grade
😊	CW1	CW2	Exam	Grade

Machine Learning settings

Quiz!

- Supervised, unsupervised or reinforcement learning?

Go to **www.menti.com**
and use the code
5635 1795



Quiz!

- Supervised, unsupervised or reinforcement learning?

“You want to identify the different distinct groups of people who are turning up at hospital A&Es with COVID.

Understanding this case-mix will help the hospital to better respond to the needs of these different types of patients.”

www.menti.com Code: 5635 1795



Quiz!

- Supervised, unsupervised or reinforcement learning?

“A book distributor has a collection of books, which it has classified into different categories, e.g. ‘Young adults’, ‘Biography’, and ‘Horror’. It wants to use this information to build a system to classify its new products automatically.”

www.menti.com Code: **5635 1795**




Quiz!

- Supervised, unsupervised or reinforcement learning?

“You work for Spotify and want to group together similar songs based on their characteristics (e.g. their tempo and their lyrics). You want to do this so you can create playlists with similar songs and suggest them to users.”

www.menti.com Code: **5635 1795**




Quiz!

- Supervised, unsupervised or reinforcement learning?

“You are trying to train a robot to open your fridge and bring you an ice-cold 0% Heineken beer.”

www.menti.com Code: 5635 1795



Classification and regression

The two most popular ML tasks

Classification

 $f($  $) = y$ \leftarrow Discrete/
categorical

Regression

 $f($  $) = y$ \leftarrow Real-valued/
continuous

Quiz!

- Classification or regression?

“You want to predict the number of people who will arrive at a hospital with COVID-19 the next day. You can use information about how many people have tested positive and historic information about how many positive tests have resulted in people arriving at hospital.”

Quiz!

- Classification or regression?

“You are trying to predict whether an individual has COVID-19 based on their current symptoms. You can use historical data from people have previously tested positive or negative and what their symptoms were at the time.”



Are you a good binary classifier?



Dataset





Samples




Dataset

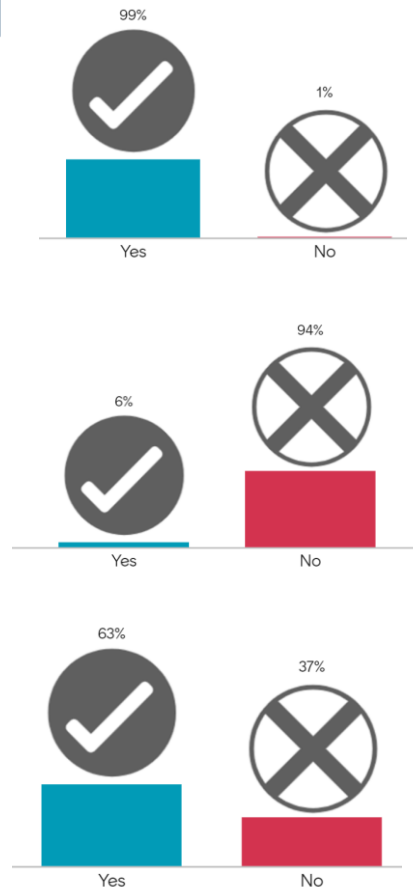


Samples






Samples



The supervised learning pipeline


Voronoi diagram




By Balu Ertl - Own work, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=38534275>

Bias-variance trade-off


One of the most important ML concepts!



Underfit



Robust, good fit



Overfit



Computational Learning Theory

Occam's razor: More things should not be used than are necessary.

Is 85% accuracy any good?

“It is all relative...”

Baseline


Chance/random performance [lower bound]

Is there a stronger baseline? e.g. most frequent class
baseline

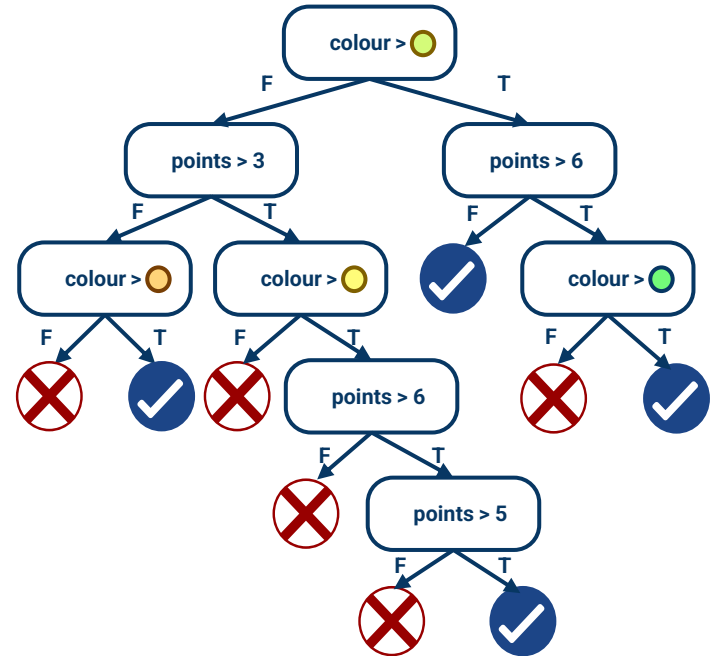
The base performance before any improvements

Coming up next week...

k Nearest Neighbours



Decision Trees



Any last questions?

Any feedback for us?

- It's anonymous!

Go to **www.menti.com**
and use the code
8244 4568



<https://www.menti.com/ns859tu54f>

THIS IS YOUR MACHINE LEARNING SYSTEM?

YUP! YOU POUR THE DATA INTO THIS BIG
PILE OF LINEAR ALGEBRA, THEN COLLECT
THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

JUST STIR THE PILE UNTIL
THEY START LOOKING RIGHT.

