


More on Gradient Descent

Yingzhen Li

Department of Computing
Imperial College London

@liyzhen2
yingzhen.li@imperial.ac.uk

October 29, 2021

Recap

Algorithm: Gradient Descent

Define **starting point** θ_0 , sequence of **step sizes** γ_t , set $t \leftarrow 0$.

1. Set $\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t)$, $t \leftarrow t + 1$
2. Repeat 1 until stopping criterion.

Recap

Algorithm: Gradient Descent **for linear regression**

Define **starting point** θ_0 , sequence of **step sizes** γ_t , set $t \leftarrow 0$.

1. Set $\theta_{t+1} = \theta_t - \gamma_t \frac{1}{\sigma^2} \mathbf{X}^\top (\mathbf{X}\theta_t - \mathbf{y})$, $t \leftarrow t + 1$
 $\Leftrightarrow \quad \theta_{t+1} = (\mathbf{I} - \frac{\gamma_t}{\sigma^2} \mathbf{X}^\top \mathbf{X}) \theta_t + \frac{\gamma_t}{\sigma^2} \mathbf{X}^\top \mathbf{y} \quad \text{if } \gamma_t = \gamma$
2. Repeat 1 until stopping criterion.

Recap

Analysing convergence of GD with constant step sizes:

$$\boldsymbol{\theta}_{t+1} = (\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{X}^\top \mathbf{X}) \boldsymbol{\theta}_t + \frac{\gamma}{\sigma^2} \mathbf{X}^\top \mathbf{y} \quad \text{if } \gamma_t = \gamma$$

Key ideas:

- ▶ Need to check if $\|\boldsymbol{\theta}_t - \boldsymbol{\theta}^*\|_2^2 \rightarrow 0$ when $t \rightarrow \infty$
- ▶ $\boldsymbol{\theta}_t - \boldsymbol{\theta}^* = \mathbf{A}^t (\boldsymbol{\theta}_0 - \boldsymbol{\theta}^*)$ where $\mathbf{A} = (\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{X}^\top \mathbf{X})$
- ▶ So we need to understand how matrix multiplications “stretch/shrink” vectors \leftarrow Eigen-decomposition & SVD
- ▶ Convergence of GD depends on $\lambda_{\max}((\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{X}^\top \mathbf{X})^2)$
 \Rightarrow need to set γ smaller than some “safe threshold”

Recap

Algorithm: Gradient Descent

Define **starting point** θ_0 , sequence of **step sizes** γ_t , set $t \leftarrow 0$.

1. Set $\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t)$, $t \leftarrow t + 1$
2. Repeat 1 until stopping criterion.

On choosing step size:


- too small: slow convergence
- too large: divergence
- just right: depends on problem (often: trial and error)

Recap


Is my choice of step size robust (to different initialisations)?

- Depending on the **condition number**:

$$\kappa(\mathbf{X}^\top \mathbf{X}) := \frac{\lambda_{\max}(\mathbf{X}^\top \mathbf{X})}{\lambda_{\min}(\mathbf{X}^\top \mathbf{X})} = \sqrt{\frac{\max_q R(\mathbf{X}^\top \mathbf{X}, q)}{\min_q R(\mathbf{X}^\top \mathbf{X}, q)}}$$



well conditioned
 $\kappa(\mathbf{X}^\top \mathbf{X}) \approx 1$



ill conditioned
 $\kappa(\mathbf{X}^\top \mathbf{X}) \gg 1$

- Need careful choice of step-sizes if the loss is “very stretched”

Gradient descent with pre-conditioning

Algorithm: Gradient Descent **with Pre-conditioning**

Define **starting point** θ_0 , sequence of **step sizes** γ_t ,

Define $\Delta\theta_t = \mathbf{0}$, set $t \leftarrow 0$.

1. **Select a pre-conditioner** \mathbf{P}_t
2. Set $\theta_{t+1} = \theta_t - \gamma_t \mathbf{P}_t^{-1} \nabla_{\theta} L(\theta_t)$
3. Set $t \leftarrow t + 1$
4. Repeat 1 - 3 until stopping criterion.

Gradient descent with pre-conditioning

Linear regression example:

Assume constant step-sizes $\gamma_t = \gamma$ and **fixed pre-conditioner \mathbf{P}** :

1. Set $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \gamma \frac{1}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top (\mathbf{X} \boldsymbol{\theta}_t - \mathbf{y})$
2. Set $t \leftarrow t + 1$
3. Repeat 1 - 2 until stopping criterion.

How to choose γ and \mathbf{P} ?

Gradient descent with pre-conditioning

Let's derive the iterative updates again:

$$\theta_{t+1} = (\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{X}) \theta_t + \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{y}$$

Gradient descent with pre-conditioning

Let's derive the iterative updates again:

$$\theta_{t+1} = (\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{X}) \theta_t + \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{y}$$

$$\Rightarrow \quad \theta_t = (\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{X})^t (\theta_0 - \theta^*) + \theta^*$$

Gradient descent with pre-conditioning

Let's derive the iterative updates again:

$$\theta_{t+1} = (\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{X}) \theta_t + \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{y}$$

$$\Rightarrow \theta_t = (\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{X})^t (\theta_0 - \theta^*) + \theta^*$$

This means we need to look into eigenvalues of $(\mathbf{I} - \frac{\gamma}{\sigma^2} \mathbf{P}^{-1} \mathbf{X}^\top \mathbf{X})^2$

1. $\lambda_{max} < 1$: always converge
2. $\lambda_{min} \geq 1$: always diverge
3. $\lambda_{min} < 1$ but $\lambda_{max} \geq 1$: depends on initialisation of θ_0

Gradient descent with pre-conditioning

To ensure convergence at **any** initialisation: $\gamma < 2\sigma^2 / \lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})$

If you want to test your luck: choose $\gamma \in [\frac{2\sigma^2}{\lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})}, \frac{2\sigma^2}{\lambda_{\min}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})})$

Is my choice of γ robust to initialisation of $\boldsymbol{\theta}_0$?

Gradient descent with pre-conditioning

To ensure convergence at **any** initialisation: $\gamma < 2\sigma^2/\lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})$

If you want to test your luck: choose $\gamma \in [\frac{2\sigma^2}{\lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})}, \frac{2\sigma^2}{\lambda_{\min}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})})$

Is my choice of γ robust to initialisation of θ_0 ?

- Depending on the **condition number**:

$$\kappa(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}) := \frac{\lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})}{\lambda_{\min}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})} = \sqrt{\frac{\max_q R(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}, q)}{\min_q R(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}, q)}}$$

Gradient descent with pre-conditioning

To ensure convergence at **any** initialisation: $\gamma < 2\sigma^2 / \lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})$


If you want to test your luck: choose $\gamma \in [\frac{2\sigma^2}{\lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})}, \frac{2\sigma^2}{\lambda_{\min}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})})$

Is my choice of γ robust to initialisation of θ_0 ?

- Depending on the **condition number**:

$$\kappa(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}) := \frac{\lambda_{\max}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})}{\lambda_{\min}(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X})} = \sqrt{\frac{\max_q R(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}, q)}{\min_q R(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}, q)}}$$

- Ideally, prefer well-conditioned optimisation: choose $\mathbf{P} \propto \mathbf{X}^\top\mathbf{X}$
- In practice: want to make $\kappa(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}) \approx 1$ while \mathbf{P}^{-1} is easy to compute



well conditioned
 $\kappa(\mathbf{P}^{-1}\mathbf{X}^\top\mathbf{X}) \approx 1$

Gradient descent with pre-conditioning

Constructing “cheap but useful” pre-conditioner \mathbf{P}_t :

- Use diagonal/low-rank matrices for \mathbf{P}_t

Gradient descent with pre-conditioning

Constructing “cheap but useful” pre-conditioner \mathbf{P}_t :

- Use diagonal/low-rank matrices for \mathbf{P}_t
- Having prior knowledge on $\nabla^2 L(\boldsymbol{\theta}_t)$ would be useful
 - For linear regression, $\nabla^2 L(\boldsymbol{\theta}_t) \propto \mathbf{X}^\top \mathbf{X}$

Gradient descent with pre-conditioning

Constructing “cheap but useful” pre-conditioner \mathbf{P}_t :

- Use diagonal/low-rank matrices for \mathbf{P}_t
- Having prior knowledge on $\nabla^2 L(\boldsymbol{\theta}_t)$ would be useful
 - For linear regression, $\nabla^2 L(\boldsymbol{\theta}_t) \propto \mathbf{X}^\top \mathbf{X}$
- $\nabla^2 L(\boldsymbol{\theta}_t)$ might not be available nor a constant
 - Approximated by statistics of gradients along the update trajectory
 - Many adaptive learning rate methods (e.g. Adam) do this!

Gradient descent with momentum

Algorithm: Gradient Descent **with Momentum**

Define **starting point** θ_0 , sequence of **step sizes** γ_t ,

Define $\Delta\theta_t = 0$, **momentum step-size** α , set $t \leftarrow 0$.

1. Set $\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t) + \alpha \Delta\theta_t$
2. Set $\Delta\theta_{t+1} = \theta_{t+1} - \theta_t = \alpha \Delta\theta_t - \gamma_t \nabla_{\theta} L(\theta_t)$
3. Set $t \leftarrow t + 1$
4. Repeat 1 - 3 until stopping criterion.

Gradient descent with momentum

$$\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t) + \alpha \Delta \theta_t$$

$$\Delta \theta_{t+1} = \theta_{t+1} - \theta_t$$

The Key idea of using momentum $\Delta \theta_t = \theta_t - \theta_{t-1}$:

Making the current iteration's update closer to the previous iterations.

Gradient descent with momentum

$$\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t) + \alpha \Delta \theta_t$$

$$\Delta \theta_{t+1} = \theta_{t+1} - \theta_t$$

The Key idea of using momentum $\Delta \theta_t = \theta_t - \theta_{t-1}$:

Making the current iteration's update closer to the previous iterations.

- Speed up in “flat regions” (i.e. $\nabla_{\theta} L(\theta_t)$ has small magnitude)

Gradient descent with momentum

$$\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t) + \alpha \Delta \theta_t$$

$$\Delta \theta_{t+1} = \theta_{t+1} - \theta_t$$

The Key idea of using momentum $\Delta \theta_t = \theta_t - \theta_{t-1}$:

Making the current iteration's update closer to the previous iterations.

- Speed up in “flat regions” (i.e. $\nabla_{\theta} L(\theta_t)$ has small magnitude)
- Alleviate oscillations (when $\nabla_{\theta} L(\theta_t)$ points to a different direction)

Gradient descent with momentum

$$\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t) + \alpha \Delta \theta_t$$

$$\Delta \theta_{t+1} = \theta_{t+1} - \theta_t$$

The Key idea of using momentum $\Delta \theta_t = \theta_t - \theta_{t-1}$:


Making the current iteration's update closer to the previous iterations.

- Speed up in “flat regions” (i.e. $\nabla_{\theta} L(\theta_t)$ has small magnitude)
- Alleviate oscillations (when $\nabla_{\theta} L(\theta_t)$ points to a different direction)
- Smooth out gradients if using inexact gradients

Momentum

Define **starting point** θ_0 , sequence of **step sizes** γ_t , set $t \leftarrow 0$.

1. Set $\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t) + \alpha(\theta_t - \theta_{t-1})$
 $t \leftarrow t + 1$
2. Repeat 1 until stopping criterion.




Without momentum

Momentum

Define **starting point** θ_0 , sequence of **step sizes** γ_t , set $t \leftarrow 0$.

1. Set $\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t) + \alpha(\theta_t - \theta_{t-1})$
 $t \leftarrow t + 1$
2. Repeat 1 until stopping criterion.



With momentum

Line search

Can we automatically adapt γ_t so we can **guarantee improvement**?

Line search

Can we automatically adapt γ_t so we can **guarantee improvement**?

Algorithm: Line Search

Define **starting point** θ_0 , sset $t \leftarrow 0$.

1. Compute gradient $\nabla_{\theta} L(\theta_t)$
2. Search $\gamma_t \in (\gamma_{min}, \gamma_{max})$ to **minimise** $L(\theta_t - \gamma \nabla_{\theta} L(\theta_t))$ **w.r.t.** γ
3. Set $\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t)$, $t \leftarrow t + 1$
4. Repeat 1-3 until stopping criterion.


Line search

Can we automatically adapt γ_t so we can **guarantee improvement**?

Algorithm: Line Search

Define **starting point** θ_0 , sset $t \leftarrow 0$.

1. Compute gradient $\nabla_{\theta} L(\theta_t)$
2. Search $\gamma_t \in (\gamma_{min}, \gamma_{max})$ to **minimise** $L(\theta_t - \gamma \nabla_{\theta} L(\theta_t))$ **w.r.t.** γ
3. Set $\theta_{t+1} = \theta_t - \gamma_t \nabla_{\theta} L(\theta_t)$, $t \leftarrow t + 1$
4. Repeat 1-3 until stopping criterion.



Line search

Algorithms for the search step:

Approximately minimise $L(\boldsymbol{\theta}_t - \gamma \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ w.r.t. $\gamma \in (\gamma_{min}, \gamma_{max})$

- Backtracking line search: shrinking step sizes

Line search

Algorithms for the search step:

Approximately minimise $L(\boldsymbol{\theta}_t - \gamma \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ w.r.t. $\gamma \in (\gamma_{min}, \gamma_{max})$

- Backtracking line search: shrinking step sizes
 - Set a “decreasing schedule” $\gamma_{max} = \gamma^{(1)} > \gamma^{(2)} > \dots > \gamma^{(\tau)} = \gamma_{min}$

Line search

Algorithms for the search step:

Approximately minimise $L(\boldsymbol{\theta}_t - \gamma \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ w.r.t. $\gamma \in (\gamma_{min}, \gamma_{max})$

- Backtracking line search: shrinking step sizes
 - Set a “decreasing schedule” $\gamma_{max} = \gamma^{(1)} > \gamma^{(2)} > \dots > \gamma^{(\tau)} = \gamma_{min}$
 - Try out these step sizes until satisfying an acceptance criterion
e.g. $L(\boldsymbol{\theta}_t) - L(\boldsymbol{\theta}_t - \gamma^{(k)} \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ is “big enough”
 - Set $\gamma_t \leftarrow \gamma^{(k)}$

Line search

Algorithms for the search step:

Approximately minimise $L(\boldsymbol{\theta}_t - \gamma \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ w.r.t. $\gamma \in (\gamma_{min}, \gamma_{max})$

- ▶ Backtracking line search: shrinking step sizes
 - ▶ Set a “decreasing schedule” $\gamma_{max} = \gamma^{(1)} > \gamma^{(2)} > \dots > \gamma^{(\tau)} = \gamma_{min}$
 - ▶ Try out these step sizes until satisfying an acceptance criterion
e.g. $L(\boldsymbol{\theta}_t) - L(\boldsymbol{\theta}_t - \gamma^{(k)} \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ is “big enough”
 - ▶ Set $\gamma_t \leftarrow \gamma^{(k)}$
- ▶ Other methods, e.g. checking the Wolfe conditions

Line search

Algorithms for the search step:

Approximately minimise $L(\boldsymbol{\theta}_t - \gamma \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ w.r.t. $\gamma \in (\gamma_{min}, \gamma_{max})$

- Backtracking line search: shrinking step sizes
 - Set a “decreasing schedule” $\gamma_{max} = \gamma^{(1)} > \gamma^{(2)} > \dots > \gamma^{(\tau)} = \gamma_{min}$
 - Try out these step sizes until satisfying an acceptance criterion
e.g. $L(\boldsymbol{\theta}_t) - L(\boldsymbol{\theta}_t - \gamma^{(k)} \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}_t))$ is “big enough”
 - Set $\gamma_t \leftarrow \gamma^{(k)}$
- Other methods, e.g. checking the Wolfe conditions

Used by e.g. nonlinear Conjugate Gradient or BFGS.

Converges **very** quickly if $\dim(\boldsymbol{\theta})$ is small.

Stochastic gradient descent

Remember objective function gradient:

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}_t} L(\boldsymbol{\theta}_t) = \nabla_{\boldsymbol{\theta}_t} \sum_{n=1}^N (f(\mathbf{x}_n; \boldsymbol{\theta}_t) - y_n)^2$$

Stochastic gradient descent

Remember objective function gradient:

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}_t} L(\boldsymbol{\theta}_t) = \sum_{n=1}^N \nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$$

Stochastic gradient descent

Remember objective function gradient:

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}_t} L(\boldsymbol{\theta}_t) = \sum_{n=1}^N \nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$$

- In big data applications N could be very large, e.g. $N > 10^9$.

Stochastic gradient descent

Remember objective function gradient:

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}_t} L(\boldsymbol{\theta}_t) = \sum_{n=1}^N \nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$$

- In big data applications N could be very large, e.g. $N > 10^9$.
- Running (full batch) gradient descent can be very expensive!
 - Need to compute $L_n(\boldsymbol{\theta}_t)$ for each data point
 - Need to store (intermediate results of) every $\nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$
 - Doing all of these only for a single update step of $\boldsymbol{\theta}$

Can we find a **computationally cheaper** but still working solution?

Stochastic gradient descent

Remember objective function gradient:

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}_t} L(\boldsymbol{\theta}_t) = \sum_{n=1}^N \nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$$

- ▶ In big data applications N could be very large, e.g. $N > 10^9$.
- ▶ **Stochastic** GD: compute a **stochastic estimator** for the sum using a **random subset**, i.e. with $M < N$ terms

$$\hat{\mathbf{g}}_t = \frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t) \quad (1)$$

Stochastic gradient descent

Remember objective function gradient:

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}_t} L(\boldsymbol{\theta}_t) = \sum_{n=1}^N \nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$$


- ▶ In big data applications N could be very large, e.g. $N > 10^9$.
- ▶ **Stochastic** GD: compute a **stochastic estimator** for the sum using a **random subset**, i.e. with $M < N$ terms

$$\hat{\mathbf{g}}_t = \frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t) \quad (1)$$

- ▶ Gradient estimator $\hat{\mathbf{g}}_t$ is now **random**!


Does optimisation still work?

Stochastic gradient descent – random gradients




- Full batch GD will converge to $\theta_t \rightarrow \theta^*$

Stochastic gradient descent – random gradients




- Full batch GD will converge to $\theta_t \rightarrow \theta^*$

Stochastic gradient descent – random gradients



- ▶ Full batch GD will converge to $\theta_t \rightarrow \theta^*$
- ▶ GD based on a single datapoint: if we select the “wrong” points, the gradient can point in the **wrong direction**!

Stochastic gradient descent – random gradients



- ▶ Full batch GD will converge to $\theta_t \rightarrow \theta^*$
- ▶ GD based on a single datapoint: if we select the “wrong” points, the gradient can point in the **wrong direction**!
- ▶ If we don't randomise, then we can keep on picking the same “wrong” point.

Stochastic gradient descent – convergence

- ▶ We choose our subsets e.g. through sampling with replacement:

$$P(\mathcal{S} = \{m_1, m_2, \dots, m_M\}) = \frac{1}{N} \cdot \frac{1}{N} \cdots = N^{-M}$$

Stochastic gradient descent – convergence

- ▶ We choose our subsets e.g. through sampling with replacement:

$$P(\mathcal{S} = \{m_1, m_2, \dots, m_M\}) = \frac{1}{N} \cdot \frac{1}{N} \cdots = N^{-M}$$

- ▶ So our estimator is **unbiased**, i.e.

$$\mathbb{E}_{\mathcal{S}}[\hat{\mathbf{g}}_t] = \mathbf{g}_t$$

Stochastic gradient descent – convergence

- ▶ We choose our subsets e.g. through sampling with replacement:

$$P(\mathcal{S} = \{m_1, m_2, \dots, m_M\}) = \frac{1}{N} \cdot \frac{1}{N} \cdots = N^{-M}$$

- ▶ So our estimator is **unbiased**, i.e.

$$\mathbb{E}_{\mathcal{S}}[\hat{\mathbf{g}}_t] = \mathbf{g}_t$$

$$\Leftrightarrow \mathbb{E}_{\mathcal{S}} \left[\frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t) \right] = \sum_{n=1}^N \nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$$

Stochastic gradient descent – convergence

Use **unbiased** gradient estimate:

$$\mathbb{E}_{\mathcal{S}} \left[\frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t) \right] = \sum_{n=1}^N \nabla_{\boldsymbol{\theta}_t} L_n(\boldsymbol{\theta}_t)$$

Stochastic gradient descent – convergence


Use **unbiased** gradient estimate:

$$\mathbb{E}_{\mathcal{S}} \left[\frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\theta_t} L_m(\theta_t) \right] = \sum_{n=1}^N \nabla_{\theta_t} L_n(\theta_t)$$

- Next, if we choose the step sizes carefully:

$$\sum_{t=1}^{\infty} \gamma_t = \infty \quad \sum_{t=1}^{\infty} \gamma_t^2 < \infty$$

- Then SGD will converge!
- See [Robbins & Monro, 1951]



Stochastic gradient descent

Choosing batch size M :

- Should be reasonably small (so can be computed fast)
- Should not be too small though:

$$P(\mathcal{S} = \{m_1, m_2, \dots, m_M\}) = \frac{1}{N} \cdot \frac{1}{N} \cdots = N^{-M},$$
$$\mathbb{V}[\hat{\mathbf{g}}_t] = \mathbb{V}_{\mathcal{S}} \left[\frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t) \right]$$

Stochastic gradient descent

Choosing batch size M :

- ▶ Should be reasonably small (so can be computed fast)
- ▶ Should not be too small though:

$$\begin{aligned}P(\mathcal{S} = \{m_1, m_2, \dots, m_M\}) &= \frac{1}{N} \cdot \frac{1}{N} \cdots = N^{-M}, \\ \mathbb{V}[\hat{\mathbf{g}}_t] &= \mathbb{V}_{\mathcal{S}} \left[\frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t) \right] \\ &= \frac{N^2}{M} \mathbb{V}_{m \sim \text{Uniform}(1, \dots, N)} [\nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t)]\end{aligned}$$

Stochastic gradient descent

Choosing batch size M :

- ▶ Should be reasonably small (so can be computed fast)
- ▶ Should not be too small though:

$$\begin{aligned}P(\mathcal{S} = \{m_1, m_2, \dots, m_M\}) &= \frac{1}{N} \cdot \frac{1}{N} \cdots = N^{-M}, \\ \mathbb{V}[\hat{\mathbf{g}}_t] &= \mathbb{V}_{\mathcal{S}} \left[\frac{N}{M} \sum_{m \in \mathcal{S}} \nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t) \right] \\ &= \frac{N^2}{M} \mathbb{V}_{m \sim \text{Uniform}(1, \dots, N)} [\nabla_{\boldsymbol{\theta}_t} L_m(\boldsymbol{\theta}_t)]\end{aligned}$$

- ▶ Advanced SGD approaches use variance reduction techniques (allowing usage of small M but still achieve small variance)

Summary

This week we discussed **gradient descent**:

- Gradient descent for linear regression
- Convergence analysis (based on eigen decomposition results)
- Choosing step sizes
- GD with advanced techniques
 - Pre-conditioning
 - Momentum
 - Line search
- Stochastic gradient descent

Summary

This week we discussed **gradient descent**:

- Gradient descent for linear regression
- Convergence analysis (based on eigen decomposition results)
- Choosing step sizes
- GD with advanced techniques
 - Pre-conditioning
 - Momentum
 - Line search
- Stochastic gradient descent

A popular optimiser for deep learning “Adam”:
SGD + momentum + pre-conditioning